

TCN, Inc. — AI Differentiation & Competitive Moat

An independent assessment of whether TCN's artificial-intelligence stack constitutes a defensible, cash-flow-protecting moat — or commodity feature parity dressed as franchise value.

PREPARED FOR

Benefit Street Partners

MANDATE

AI / technology advisory

SECTOR

CCaaS · Collections / ARM software

DATE

June 2026

PREPARED BY

K. Slowinski
Independent AI & Technology Advisor

Basis of preparation

Scope & purpose

This material was prepared at the request of Benefit Street Partners ("BSP") to support internal evaluation of TCN, Inc. ("TCN" or the "Company"). It addresses AI/technology differentiation and competitive moat only, and is not a complete business, financial, legal, or tax assessment.

Sources & estimates

Analysis relies on public company materials, competitor disclosures, third-party review corpora, and primary technical reconstruction. Financial and headcount figures are third-party estimates, are unaudited, and are labelled where used. No representation or warranty, express or implied, is given as to accuracy or completeness.

No advice / forward-looking

Nothing herein is investment, legal, or regulatory advice or a recommendation to transact. Forward-looking statements involve risk and may differ materially from outcomes. Recipients should perform independent diligence.

Confidentiality

For the internal use of BSP investment and diligence teams only. Not to be reproduced or distributed, in whole or in part, without prior written consent of the author.

The headline: a vertical franchise, not an AI franchise

VERDICT

TCN has no durable AI moat. Its AI is at competitive parity, built largely on commodity third-party models, and replicable in days. We could not identify a single AI capability TCN offers that its major competitors do not.

- **Parity, not edge.** Agent assist, summaries, analytics, voicebots, auto-QA — all table stakes also shipped by Five9, NICE and Talkdesk, often more maturely.
- **Commodity stack.** Speech, LLMs and even compliance data are bought, not built. TCN is an integrator, not a model owner.
- **The real moat is elsewhere** — vertical/compliance depth in collections and switching costs — and it is only **modest**.
- **The only AI-native moat is labor substitution.** Replacing human agents with autonomous AI agents. No vendor — TCN included — has done it at scale. Until then, AI is a feature, not a franchise.

RECOMMENDATION — CREDIT LENS

*Underwrite to the durable vertical + switching-cost cash flows. Do **not** capitalize an AI premium. Treat agentic "agent replacement" as unfunded optionality, not base case.*

What would make AI a moat — and how we tested it

A capability is only a moat if a well-funded competitor cannot cheaply replicate it or take the customer. We scored TCN against five tests.

Test 1

Proprietary models / IP

Own models, training pipeline, or technical IP rivals lack.

Test 2

Proprietary data & network effects

Unique data that compounds and improves the product with scale.

Test 3

Workflow lock-in

Switching costs, integrations, data gravity, audit/legal record.

Test 4

Distribution / vertical depth

Channel, brand, and domain depth in a defensible niche.

Test 5 — the decisive one

Economic step-change: labor substitution

Does the AI remove the customer's largest cost — human agents — rather than merely assist them? This is the only AI capability that re-bases unit economics and creates new lock-in.

We weight it highest.

Methodology: public materials & competitor disclosures · review corpora (G2/Capterra) · third-party headcount/financial estimates · hands-on reconstruction of TCN's AI components.

TCN, Inc. — small, profitable, vertically focused

ARR (2025A)

~\$60M

~\$69M Q1'26 · ~37% 2026E

Customers · ARPA

1,000+

~\$60K ARPA (up from \$42K)

Profitability

~24%

Adj. EBITDA · 75%+ GM · Rule-of-60

Retention

112%

NRR · 81% GRR · 98% logo

Profile

Founded 1999, St. George, Utah. Flagship platform **TCN Operator** (six modules: omnichannel, compliance/data, analytics, workforce, automation, infrastructure). Verticals: **collections / accounts-receivable management**, BPO, healthcare RCM, utilities; actively expanding UK/EU.

Credit read

A genuinely attractive niche credit: ~\$60M ARR, ~30%+ organic growth, ~24% EBITDA, bootstrapped (**\$0 raised since 1999**). But still **15–45× smaller** than the CCaaS leaders — **sub-scale to win an AI R&D arms race**. Underwrite the franchise, not the AI narrative.

The AI toolbox — and what's actually under the hood

COMPONENT	MARKETED AS	UNDER THE HOOD (OUR ASSESSMENT)
Conversational / Voice Analytics	Transcribe, search, flag 100% of calls	Third-party ASR + keyword/intent scoring — commodity
Agent Assist	Real-time coaching, next-best-action, summaries	LLM layer over transcripts — commodity wrapper
Natural Language Compliance (NLC)	"Industry-first" plain-English compliance rules	LLM English → rule parser on a rules engine; value is the audit trail, not the AI
Chatbots / Email automation	Deflect & automate routine inbound	LLM + RAG — commodity
Voicebot / AI IVR	Self-service voice, intent routing	ASR + LLM + TTS pipeline — commodity
Auto-QA / Workforce Opt.	Auto-score every conversation	Scoring rubric over transcripts — table stakes

SIGNAL

Engineering job posts cite PostgreSQL + ClickHouse and "working closely with state-of-the-art AI models." TCN *consumes* models; there is no evidence of a proprietary model, training pipeline, or ML research function.

Nothing in the toolbox a major peer doesn't already have

AI CAPABILITY	TCN	FIVE9 (GENIUS)	NICE (ENLIGHTEN/MPOWER)	TALKDESK (ASCEND)	TCN EDGE?
Real-time agent assist / copilot	●	●		●	None
Automated call summaries	●	●		●	None
Conversational / voice analytics	●	●		●	None
Self-service voicebot / AI IVR	●	●		●	None
Auto-QA (100% of calls)	●	●		●	None
No-code AI agent builder	○	●		●	Behind
Autonomous / agentic AI (acts end-to-end)	◐	●		●	Behind
Proprietary models / data moat	○	◐		◐	None
Compliance / TCPA depth (collections)	●	◐		◐	Vertical, not AI

● shipping ◐ partial / emerging ○ absent or trailing. TCN's only relative strength is collections-compliance depth — a **domain** advantage, not an AI one. On true agentic AI it trails the leaders.

A niche specialist in a field of giants

VENDOR	EST. REVENUE / ARR	STRUCTURE	AI SUITE	STRATEGIC FOCUS
NICE	~\$2.7B	Public	Enlighten / Mpower	Enterprise CX, WFO, compliance
Genesys	~\$2.3B	Private	Cloud CX AI	Enterprise omnichannel at scale
Five9	~\$1.0B	Public	Genius AI	Outbound / blended CX, SMB–ent.
Talkdesk	~\$0.2B	Private	Ascend AI	Mid-market / enterprise, AI-first
TCN	~\$60M	Bootstrapped	"Operator" AI	Collections / ARM vertical



IMPLICATION

TCN is roughly 15–45× **smaller** than the category leaders. It cannot out-spend them on AI R&D; it can only win by **vertical depth and price**. That is a defensible niche — not an AI arms-race position.

Compliance is table stakes, not a moat

The CIP and the commercial VDD both lean on a "compliance moat." But you cannot sell into collections without compliance — and every competitor, including the AI-native disruptors, markets it.

VENDOR	MARKETS COLLECTIONS COMPLIANCE?	WHAT THEY CLAIM
LiveVox (now NICE)	• Yes — its heritage	Pioneered "Four Clouds" / Human Call Initiator; court-validated non-ATDS
Five9	• Yes	Dedicated TCPA and debt-collections product pages
Convoso	• Yes	Compliance product: dialing guardrails + automation in-workflow
Skit.ai (AI-native)	• Yes — headline	"Automatically follows FDCPA, TCPA, Reg F" — TCN's NLC pitch, verbatim
NICE / Genesys / Talkdesk	• Yes	Compliance tooling, audit trails, Reg-F templates
DNC.com / Gryphon / PossibleNOW	• Yes	Compliance-as-a-service — buyable a la carte

WHY IT MATTERS

The AI-native entrants are **compliance-native too** — so compliance does not shield TCN from disintermediation. Strip the label and the "compliance moat" is really **switching costs + being the last independent mid-market specialist** after LiveVox was absorbed by NICE. A market position, not a technical moat.

The AI is rented; the moat would have to be built

Bought / commodity

- ASR — Whisper / Deepgram / AWS Transcribe class
- LLMs — OpenAI / Anthropic / AWS
- TTS — ElevenLabs / Cartesia / Polly class
- **Compliance data** — DNC, litigator & reassigned-number lists are a commodity market (DNC.com: 70B+ scrubs; PossibleNOW; Gryphon; Blacklist Alliance)


Built / defensible

- Integrated platform & dialer
- Contact-eligibility rules engine + NL authoring
- **Consent / revocation audit ledger** (the courtroom record)
- Collections-specific workflows & data gravity

REPRODUCIBILITY TEST — THE IP BARRIER IS LOW

*As part of this diligence we rebuilt TCN's four headline AI components **plus** a tamper-evident consent/audit ledger over a 100-record dataset — offline, with no proprietary data, in a single afternoon. A capability that reproduces that quickly is not a moat.*

The castle is the vertical. The AI is paint on the drawbridge.



What's real

Purpose-built for debt collection (a customer literally cites "solutions designed for the debt-collection industry"), reinforced by switching costs and the legal-defensibility of consent/audit records.

What's not

The AI. It neither creates lock-in beyond the platform itself nor compounds with proprietary data. Removing the "AI" branding would not change TCN's competitive position.

The only AI moat is labor substitution

Contact-center cost stack

~60–70%

is human agent labor — the prize

Augmentation (today)

Copilots & summaries trim handle time. **Incremental** — and equally available to every competitor. No lock-in.

Substitution (the moat)

Replace seats with autonomous agents → re-bases unit economics, shifts pricing seat → outcome, and creates an automated-resolution **data flywheel** + deep lock-in.

THESIS

*A moat is created only by convincing customers to remove human agents and replace them with AI agents. Industry-wide that conversion has **not** happened at scale — "agentic" is still largely demoware. The first mover to convert seats into agents captures the labor budget and the lock-in. **Until TCN does this, there is no moat.***

TCN-specific angle

Collections is unusually automatable — scripted, repetitive, compliance-bounded — so the **opportunity is genuine**. But TCN sells augmentation today and is sub-scale on R&D. This is the single most important item to diligence and to fund.

Scored against the five tests

MOAT TEST	RATING	RATIONALE
1 · Proprietary models / IP	LOW	No owned models or training pipeline; consumes third-party AI.
2 · Proprietary data / network effects	LOW	Compliance data is bought; no compounding data advantage.
3 · Workflow lock-in / switching costs	MEDIUM	Real via integration + consent/audit records; not AI-driven.
4 · Distribution / vertical depth	MED-HIGH	Genuine collections/ARM domain depth; efficient GTM.
5 · Labor substitution (weighted x)	NOT YET	Sells augmentation, not replacement. The decisive, unrealized lever.
Overall AI moat	LOW	Niche franchise with modest, non-AI defensibility. No AI premium warranted.

Ratings are qualitative judgments for diligence framing, not a quantitative score.

KPIs & catalysts to monitor

Substitution is happening (bull signals)

- % interactions **fully automated** / containment rate rising
- Reference customers **removing agent seats**
- Pricing shifts from **per-seat** → **per-outcome**
- **Gross margin expansion** from automation
- Agentic product GA with named logos

Still a feature (bear signals)

- AI sold only as add-on; **low SKU attach / ARPU uplift**
- **R&D % of revenue** trails peers (scale gap widens)
- Logo churn in collections vertical
- Larger peers bundle **agentic + compliance** and undercut
- Net revenue retention flat/declining

DILIGENCE ASK

Request the automation-rate cohort data and the AI-SKU attach/ARPU bridge. These two series determine whether the AI optionality is real — or narrative.

Turning the optionality into a moat

The agent-replacement opportunity is real but unfunded. BSP capital is the unlock. A staged path from augmentation to substitution:

Horizon 1 · 0–12 mo

Instrument & harden

Productize the consent/audit ledger as defensible IP. Add automation-rate telemetry. Pick 2–3 high-volume collections workflows to fully automate (balance/payment, promise-to-pay, first-party reminders).

Horizon 2 · 12–24 mo

Substitute & re-price

Ship autonomous voice agents for those workflows. Shift pricing **seat** → **outcome** on automated volume. Publish containment & right-party metrics with reference logos.

Horizon 3 · 24–36 mo

Re-base economics

Automated resolutions become the majority of volume. Data flywheel compounds. Extend into healthcare RCM. Defend with the compliance + audit moat.

VALUE BRIDGE

*Success shows up as rising **containment %**, **gross-margin expansion**, **AI ARPU uplift** and **NRR** — the metrics that would re-rate the asset. Absent execution, none of this is in the base case.*

Underwrite the franchise, not the narrative

- **Cash-flow durability** rests on the vertical + switching costs — underwrite leverage and value there, not on AI.
- **Valuation discipline:** resist any AI multiple premium. Agentic upside = **optionality**, not base case.
- **Downside is contained** by sticky collections workflows and audit-record gravity — supportive of a credit thesis.
- **Value-creation if invested:** (1) lead agentic **substitution** in collections; (2) productize the compliance/audit ledger as the moat; (3) shift pricing seat → outcome.
- **Key risks:** sub-scale R&D, CCaaS commoditization, peers bundling agentic + compliance, regulatory shifts (TCPA/Reg F).

BOTTOM LINE

*A solid niche credit with a modest, vertical/switching-cost moat — and **no AI moat today**. The AI optionality becomes real only if management commits to replacing human agents with AI agents. Price the base case; treat substitution as the upside call option, not the thesis.*

Sources, method & glossary

Primary sources

TCN (tcn.com): AI, Operator, NLC, voice-analytics, careers, case studies, press releases (2025).
 Competitors: Five9, NICE, Talkdesk newsrooms; CX Today.
 Compliance market: DNC.com / Contact Center Compliance, PossibleNOW, Blacklist Alliance.
 Profiles: Crunchbase, ZoomInfo, PitchBook, Top Workplaces; G2 / Capterra reviews.
 BSP: Franklin Templeton disclosures (2025–26).

Method

Triangulation of public materials, competitor disclosures and review corpora; third-party financial/headcount estimates (unaudited); and a primary technical reconstruction of TCN's AI components incl. a tamper-evident consent/audit ledger over 100 records. Ratings are qualitative diligence framing.

Glossary

CCaaS — contact center as a service. **ARM** — accounts-receivable management (collections). **ASR** — automatic speech recognition. **NLC** — TCN's Natural Language Compliance. **Agentic AI** — AI that autonomously completes multi-step tasks. **TCPA / Reg F** — US telemarketing & debt-collection rules.

Companion artifact: a runnable reconstruction of TCN's AI stack (four components + audit ledger) accompanies this deck and is available on request.