

TCN AI Differentiation & Competitive Moat

An independent technology assessment of TCN's AI, prepared for BSP's credit evaluation.

PREPARED FOR

Benefit Street Partners

MANDATE

AI / technology advisory

SECTOR

CCaaS · Collections / ARM software

DATE

June 2026

PREPARED BY

K. Slowinski
Independent AI & Technology Advisor

Basis of preparation

Scope & purpose

This material was prepared at the request of Benefit Street Partners ("BSP") to support internal evaluation of TCN ("TCN" or the "Company"). It addresses AI/technology differentiation and competitive moat only, and is not a complete business, financial, legal, or tax assessment.

Sources & scope

Analysis relies on the Company's CIP and the PASG commercial VDD (provided through the process), public company and competitor materials, third-party estimates (labelled), and a primary technical reconstruction of the AI components. It does NOT include customer interviews, data-room or code access, or management discussions (a management call is the recommended next step). Figures are unaudited; no representation or warranty is given.

No advice / forward-looking

Nothing herein is investment, legal, or regulatory advice or a recommendation to transact. Forward-looking statements involve risk and may differ materially from outcomes. Recipients should perform independent diligence.

Confidentiality

For the internal use of BSP investment and diligence teams only. Not to be reproduced or distributed, in whole or in part, without prior written consent of the author.

The headline: good calling business, no AI moat

VERDICT

TCN has no durable AI moat. The generative layer runs on rented third-party models, and there was nothing we could find that clearly demonstrates an AI moat; TCN's real edges are vertical and workflow, not AI. The AI story is a tailwind only if the call stays on TCN's line. That condition is the whole credit question.

- **No edge.** Agent assist, summaries, analytics, voicebots, auto-QA. All table stakes, and Five9, NICE and Talkdesk ship them too, often more maturely.
- **Rented stack.** Speech and generative AI are third-party (GPT, Gemini, Mistral, Whisper); compliance data is buyable. One real in-house model: SmartAMD.
- **The real moat is elsewhere.** Vertical/compliance depth in collections and switching costs. And it is only **modest**.
- **Behind, not ahead.** The only AI-native moat is labor substitution. And TCN's agentic AI had **0 customers and \$0 ARR at end-2025** (first customer Apr 2026). AI-native players are built differently. Until TCN converts seats to agents, AI is a feature, not a franchise.

RECOMMENDATION. CREDIT LENS

*It's a good calling business. Underwrite the loan on that: switching costs, a sticky-ish base, and **consumption pricing** (AI volume makes them money instead of killing seats). Don't credit the AI story, and treat being behind on AI as a headwind to the credit.*

What would make AI a moat. And how we tested it

A capability is only a moat if a well-funded competitor cannot cheaply replicate it or take the customer. We scored TCN against five tests.

Test 1

Proprietary models / IP

Own models, training pipeline, or technical IP rivals lack.

Test 2

Proprietary data & network effects

Unique data that compounds and improves the product with scale.

Test 3

Workflow lock-in

Switching costs, integrations, data gravity, audit/legal record.

Test 4

Distribution / vertical depth

Channel, brand, and domain depth in a defensible niche.

Test 5. The decisive one

Economic step-change: labor substitution

Does the AI remove the customer's largest cost. Human agents. Rather than merely assist them? This is the only AI capability that re-bases unit economics and creates new lock-in. **We weight it highest.**

Methodology: public materials & competitor disclosures · review corpora (G2/Capterra) · third-party headcount/financial estimates · hands-on reconstruction of TCN's AI components.

TCN: small, profitable, vertically focused

ARR (2025A)

~\$60M

~\$69M Q1'26 · ~37% 2026E

Customers · ARPA

~1,200

~\$60K ARPA (up from \$42K)

Profitability

~20→24%

EBITDA LTM → 2026E, 75% GM, NPS 52

Retention

112%

NRR · 81% GRR · 98% logo

Profile


Founded 1999, St. George, Utah. Flagship platform **TCN Operator** (six modules: omnichannel, compliance/data, analytics, workforce, automation, infrastructure). Verticals: **collections / accounts-receivable management**, BPO, healthcare RCM, utilities; actively expanding UK/EU.

Credit read

A genuinely attractive niche credit: ~\$60M ARR, ~37% 2026E growth, ~20 → 24% EBITDA, clears Rule of 40 (their "Rule of 60" is a 2026E projection), bootstrapped (**\$0 raised since 1999**). But **~14-39x smaller** than the three scaled leaders (Five9/Genesys/NICE, on \$69M Mar-26 run-rate). Too small to win an AI R&D arms race. Underwrite the calling business, not the AI story.

A niche specialist in a field of giants

VENDOR	EST. REVENUE / ARR	STRUCTURE	AI SUITE	STRATEGIC FOCUS
NICE	~\$2.7B	Public	Enlighten / Mpower	Enterprise CX, WFO, compliance
Genesys	~\$2.3B	Private	Cloud CX AI	Enterprise omnichannel at scale
Five9	~\$1.0B	Public	Genius AI	Outbound / blended CX, SMB-ent.
Talkdesk	~\$0.2B	Private	Ascend AI	Mid-market / enterprise, AI-first
TCN	~\$69M	Bootstrapped	"Operator" AI	Collections / ARM vertical



IMPLICATION

TCN is ~14-39× **smaller** than the three scaled leaders (Five9/Genesys/NICE, vs \$69M Mar-26 run-rate); only AI-first Talkdesk (~\$0.2B) is within ~3×. It can't out-spend the majors on AI R&D. It wins by **vertical depth and price**. A defensible niche, not an arms-race position.

The AI toolbox. And what's actually under the hood

COMPONENT	MARKETED AS	UNDER THE HOOD (OUR ASSESSMENT)
Conversational / Voice Analytics	Transcribe, search, flag 100% of calls	Third-party ASR + keyword/intent scoring. commodity
Agent Assist	Real-time coaching, next-best-action, summaries	LLM layer over transcripts. commodity wrapper
Natural Language Compliance (NLC)	"Industry-first" plain-English compliance rules	LLM English → rule parser on a rules engine; value is the audit trail, not the AI
Chatbots / Email automation	Deflect & automate routine inbound	LLM + RAG. commodity
Voicebot / AI IVR	Self-service voice, intent routing	ASR + LLM + TTS pipeline. commodity
Auto-QA / Workforce Opt.	Auto-score every conversation	Scoring rubric over transcripts. table stakes

SIGNAL

*TCN's generative AI (transcription, summaries, QA) runs on third-party LLMs (GPT, Gemini, Mistral, Whisper; CIP p. 36). Its one in-production proprietary model, **SmartAMD**, is narrow (call classification) but real: answering-machine detection is now an adversarial problem (iOS, Pixel and carrier screeners sit between dialers and humans), and TCN retrains it on its own outcome-labeled dial data at scale. That's a genuine, compounding data loop, and a reason calls stay on TCN's line. It is not a generative or agentic moat. (The CIP's "majority of AI runs on TCN's native stack" counts per-call SmartAMD executions; every generative feature is third-party.)*

Nothing in the toolbox a major peer doesn't already have

AI CAPABILITY	TCN	FIVE9 (GENIUS)	NICE (ENLIGHTEN/MPOWER)	TALKDESK (ASCEND)	TCN EDGE?
Real-time agent assist / copilot	●	●	●	●	None
Automated call summaries	●	●	●	●	None
Conversational / voice analytics	●	●	●	●	None
Self-service voicebot / AI IVR	●	●	●	●	None
Auto-QA (100% of calls)	●	●	●	●	None
No-code AI agent builder	○	●	●	●	Behind
Autonomous / agentic AI (acts end-to-end)	◐ (a)	◐ (b)	◐ (b)	◐ (b)	Behind on proof
Proprietary models / data moat	◐	◐	●	◐	None
Compliance / TCPA depth (collections)	●	◐	●	◐	Vertical, not AI

● shipping ◐ partial / emerging ○ absent or trailing. TCN's only relative strength is collections-compliance depth. A **domain** advantage, not an AI one. Its only in-production proprietary model (SmartAMD) is narrow call classification. (a) TCN: first agentic customer Apr-2026, immaterial ARR. (b) Peers: shipping but early, mostly inbound self-service; nobody has scaled autonomous collections. Capability claims per vendor materials, not verified.

Compliance is table stakes, not a moat

The CIP and the commercial VDD both lean on a "compliance moat." But you cannot sell into collections without compliance. And every competitor, including the AI-native disruptors, markets it.

VENDOR	MARKETS COLLECTIONS COMPLIANCE?	WHAT THEY MARKET (CLAIMS, NOT VERIFIED)
LiveVox (now NICE)	● Yes. Its heritage	Pioneered "Four Clouds" / Human Call Initiator; court-validated non-ATDS
Five9	● Yes	Dedicated TCPA and debt-collections product pages
Convoso	● Yes	Compliance product: dialing guardrails + automation in-workflow
Skit.ai (AI-native)	● Yes. Headline	"Automatically follows FDCPA, TCPA, Reg F". TCN's NLC pitch, verbatim
NICE / Genesys / Talkdesk	● Yes	Compliance tooling, audit trails, Reg-F templates
DNC.com / Gryphon / PossibleNOW	● Yes	Compliance-as-a-service. Buyable a la carte

WHY IT MATTERS

The VDD's interviews do say peers "lack the compliance features" for collections at scale, and we credit TCN's depth. But compliance marketing is table stakes, the data underneath is buyable, and the AI-natives are shipping compliance-native. Depth is a modest, real **vertical/legal** advantage. Not an AI one. Strip the label and the "moat" is switching costs + being the last independent specialist after LiveVox was absorbed by NICE.

The AI is rented; the moat would have to be built

Bought / commodity

- ASR. Whisper / Deepgram / AWS Transcribe class
- LLMs. OpenAI / Anthropic / AWS
- TTS. ElevenLabs / Cartesia / Polly class
- **Compliance data.** DNC, litigator & reassigned-number lists are a commodity market (DNC.com: 70B+ scrubs; PossibleNOW; Gryphon; Blacklist Alliance)


Built / defensible

- Integrated platform & dialer
- Contact-eligibility rules engine + NL authoring
- **Consent / revocation audit ledger** (the courtroom record)
- Collections-specific workflows & data gravity

REPRODUCIBILITY TEST. THE IP BARRIER IS LOW

*As part of this diligence we reproduced TCN's four headline **AI/LLM components** offline over a 100-record dataset in an afternoon. Confirming the AI layer is commodity. What we did **not** reproduce is the production consent/audit ledger at customer scale, with years of legally-defensible history. Which is exactly where the modest, non-AI moat sits.*

Where the real defensibility is (it isn't AI)



What's real

Purpose-built for debt collection (a customer literally cites "solutions designed for the debt-collection industry"), reinforced by switching costs and the legal-defensibility of consent/audit records.

What's not

The AI. It neither creates lock-in beyond the platform itself nor compounds with proprietary data. Removing the "AI" branding would not change TCN's competitive position.

The only AI moat is labor substitution

Contact-center cost stack

~60-70%

is human agent labor. The prize

Augmentation (today)

Copilots & summaries trim handle time. **Incremental.**
And equally available to every competitor. No lock-in.

Substitution (the moat)

Replace seats with autonomous agents → re-bases unit economics, shifts pricing seat → outcome, and creates an automated-resolution **data flywheel** + deep lock-in.

THESIS

*A moat is created only by convincing customers to remove human agents and replace them with AI agents. Industry-wide that conversion has **not** happened at scale. "agentic" is still early everywhere. The first mover to convert seats into agents captures the labor budget and the lock-in. **Until TCN does this, there is no moat.***

Where TCN stands

Everything shipped so far is AI **for the human**: assist, summaries, QA, SmartAMD. Useful, but a person is still on every call, and no labor dollars have moved to TCN yet. Collections is automatable in layers: first-party reminders and payment-taking (high containment plausible) vs third-party negotiation (low), which is where the labor dollars actually sit. Late, but closable: the durable risks are per-outcome pricing disruption and losing the system-of-record position, not model capability.

The team question

The CIP claims a "proven global R&D engine," "120+ engineers" and "in-house model training." What it doesn't name is any AI/ML leadership or senior AI hire. Leadership is the founding Bird family plus long-tenured contact-center engineers (public bios; see appendix). Credible dialer team; unproven AI org. Test this on the call.

Who's coming for the call volume, and what they ship today

COMPANY	WHAT THEIR AI DOES	END TO END?	OWNS THE LINE? (1)	CAPITAL	THREAT TO TCN
TCN	Assists the human agent (summaries, QA, SmartAMD); AI Voice Agent launched Apr-26	◀ first customer Apr-26	● its dialer + numbers	\$0 raised (bootstrapped)	n/a. The incumbent
Five9 / NICE	Copilots plus early agentic, mostly inbound self-service	◀ early	● their dialers	Public, \$1-2.7B rev	Bundling agents + compliance above TCN
Decagon	Autonomous agents resolving chat/email/voice end to end; auditable plain-English compliance rules	● yes	◀ entering voice	~\$400M raised, \$4.5B val	Highest. Outbound voice + compliance, well-funded
Sierra	Autonomous voice + digital agents, outcome pricing; Bret Taylor	● yes	◀ voice via SIP	~\$1B+ raised, \$15.8B val	High. Capital + trust to enter regulated
Skit.ai	Collections-native: runs the whole outbound collections call with FDCPA/TCPA/Reg F built in	● yes	◀ bundled telephony (2)	~\$48M raised	Most direct. TCN's exact use case and buyer
Cresta / Parloa	Agent assist + voice-first agentic at enterprise scale	◀ partial	○	\$1.6B / \$3.0B val	Overlaps assist, analytics, voice automation
Intercom (Fin)	Fin Voice: real-time multi-step phone workflows, can call outbound	● yes	◀ rides customer SIP	Public	Voice live and expanding from support CX

BOTTOM LINE

Holds for TCN: consumption pricing, dial-path incumbency, number reputation (STIR/SHAKEN attestation, spam-label remediation), compliance depth, switching costs. **Eroding:** the disruptors are compliance-native, several run calls end to end today, and their per-minute economics work. Their

Autonomous collections calls are a legal minefield, for everyone

The rules that bind an AI caller

- FCC (Feb 2024): AI-generated voices are "artificial voice" under the **TCPA**. Prior express consent required for cell phones; \$500-\$1,500 statutory damages per call, class-action fuel.
- **FDCPA §807**: a hallucinated balance or misstated settlement is a false representation. Every AI call is recorded, discoverable evidence.
- **Reg F**: 7-in-7 attempt caps counted per debtor across channels; the attempt ledger must be airtight.
- State AI laws landing: bot-disclosure (CA), Utah AI Policy Act (TCN's home state), Colorado AI Act (2026).

Why this cuts both ways

It slows the AI-natives more in collections than in support CX: whoever holds the consent, revocation and attempt ledger controls TCPA/Reg-F compliance for any AI agent placed on top. That is TCN's system-of-record advantage, and it gets **more** valuable in the AI era.

It's also execution risk for TCN's own agent plan: autonomous negotiation with near-zero error tolerance, constrained generation for amounts and disclosures, human escalation triggers, and 100% post-call audit. None of that is demonstrated yet.

CREDIT READ

Regulation is the strongest brake on the displacement scenario and the strongest reason TCN's consent/audit ledger keeps value through the loan term. It is also the reason to discount everyone's agentic revenue plans in this vertical, TCN's included.

They measure a better dialer, not better collections

TCN can quantify a better dialer and cheaper QA. It cannot yet quantify the metrics that actually decide collections. The ones the AI-natives lead with.

METRIC THAT MATTERS TO A COLLECTIONS STAKEHOLDER	WHAT TCN CAN SHOW TODAY	THE AI-NATIVE BAR (SKIT.AI / PRODIGAL)
Right-party contact / connect	RPC 7.68% → 8.70%, +13% agent utilization, 8.8× ROI (SmartAMD)	RPC tracked across 1B+ conversations, \$1B+ resolved
QA / compliance coverage	100% of calls AI-scored, 50% fewer false positives, 30% infra cut	FDCPA/TCPA/Reg F enforced natively in the agent
Autonomous resolution / containment %	Zero through end-2025: 0 customers / \$0 (first customer Apr 2026; CIP pp. 54-55)	AI handles high-volume follow-ups; escalates only complex cases
Cost per resolution vs. human	Not reported	Tracked per client; 10× ROI case studies
Liquidation / recovery lift from AI	Not reported (connect-rate only)	Intent-to-pay scoring lifts liquidation, cuts collection time (Prodigal)

BOTTOM LINE

*Their numbers prove a good dialer and cheaper QA. That's efficiency. This category gets won on **containment, cost-per-resolution and liquidation**, and TCN can't report one of those yet. Strip the AI label and the numbers say: a very good calling platform.*

The margin is plausible today. The flat price is not

What an AI voice minute costs (mid-2026)

- Speech-to-text: ~\$0.005-0.01
- LLM inference: ~\$0.01-0.03
- Text-to-speech: ~\$0.02-0.04
- Telephony: ~\$0.005-0.01

All-in COGS: ~\$0.05-0.10/min. At the CIP's \$0.30/min price, TCN's claimed ~80% AI gross margin is credible *today*.

The assumption that breaks

The CIP's model holds the AI price at **\$0.30/min flat for five years** to reach **\$174M+ of AI agent ARR** (CIP pp. 68-69). Meanwhile model costs have been falling ~40%+ per year, self-serve voice-agent platforms retail at \$0.07-0.15/min, and the AI-natives price per **outcome**, not per minute. Price compression, not technical feasibility, is what kills that curve.

CREDIT READ

Don't credit the \$174M. The AI upside case rests on a five-year flat price in the fastest-deflating pricing environment in software. Feasible product, fantasy price curve.

What displacement does to the revenue that services the debt

Q1'26 product mix (CIP p. 54): Outbound \$41M (60%), Inbound \$13M (19%), SMS \$5M, BI \$4M, other ~\$6M. Outbound minutes are the displaceable core; compliance, recording and number management are the sticky tail.

SCENARIO OVER THE LOAN TERM	OUTBOUND ARR IMPACT	TOTAL ARR IMPACT (~\$69M BASE)	NOTE
AI-natives take 10% of outbound minutes	-\$4M	~ -6%	Absorbable; watch the trend
AI-natives take 25% of outbound minutes	-\$10M	~ -15%	Material to coverage; GRR already shows downsell
AI-natives take 50% of outbound minutes	-\$21M	~ -30%	Thesis-breaking for the credit
Upside branch: minutes migrate to TCN's own AI	up to ~6x per migrated minute	net positive	Only if the call stays on TCN's line

OPEN DILIGENCE ITEM

Illustrative static math, not a model: it assumes displaced minutes are lost at current per-minute pricing and ignores volume growth. The precise inputs a lender needs: committed-minimum vs volume-variable ARR split, and the GRR bridge (churn vs downsell). Both are on the management-call list.

The bull case, steelmanned, and what would prove us wrong

STRONGEST COUNTER-EVIDENCE (PASG VDD, 31 INTERVIEWS + 80-RESPONDENT SURVEY)	WHY IT ISN'T DISPOSITIVE
73% of TCN customers prefer to buy AI from their existing platform vendor	Preference for the incumbent, surveyed before AI-natives had reference deployments in collections. Preferences follow proof; ask again after a Skit.ai win in the base.
"AI Agents make calls via the same dialer," so AI is limited disruption risk	True only while the dialer is the default path. The scenario to test is the agent that brings its own line, which now exists.
Switching a dialer is hard; downtime is revenue risk	Real, and we credit it (switching costs are half our moat case). But agent adoption isn't a dialer switch; it can start on one campaign without ripping anything out.
SmartAMD's 8.8x ROI shows real, compounding AI	Agreed, and we credit it as TCN's one real data loop. It's narrow (call classification), and it defends the line rather than wins the agent race.

WHAT WOULD CHANGE OUR VIEW

A working live demo of an end-to-end autonomous collections call on TCN's line; real containment and cost-per-resolution data from the April-2026 customer; contractual or technical dial-path lock-in; named senior AI/ML hires; and a dated, funded catch-up roadmap. Any three of those and "behind" becomes "closing the gap."

Scored against the five tests

MOAT TEST	RATING	RATIONALE
1 · Proprietary models / IP	LOW	One narrow proprietary model (SmartAMD, AMD); generative AI is third-party. No franchise-grade IP.
2 · Proprietary data / network effects	LOW	Compliance data is bought. One exception: SmartAMD's outcome-labeled dial data compounds, but it's narrow.
3 · Workflow lock-in / switching costs	MEDIUM	Real via integration + consent/audit records; not AI-driven.
4 · Distribution / vertical depth	MED-HIGH	Genuine collections/ARM domain depth; efficient GTM.
5 · Labor substitution (weighted x)	NOT YET	Sells augmentation, not replacement. The decisive, unrealized lever.
Overall AI moat	LOW	Niche franchise with modest, non-AI defensibility. No AI premium warranted.

Ratings are qualitative judgments for diligence framing, not a quantitative score.

KPIs & catalysts to monitor

Substitution is happening (bull signals)

- % **interactions fully automated** / containment rate rising
- Reference customers **removing agent seats**
- Pricing shifts from **per-seat** → **per-outcome**
- **Gross margin expansion** from automation
- Agentic product GA with named logos

Still a feature (bear signals)

- AI sold only as add-on; **low SKU attach / ARPU uplift**
- **R&D % of revenue** trails peers (scale gap widens)
- Logo churn in collections vertical
- Larger peers bundle **agentic + compliance** and undercut
- Net revenue retention flat/declining
- Forward-deployed cost per agent deployment not declining across cohorts (FDE itself is the winning GTM; static FDE cost is the bear case)

Quarterly reporting we'd require as lender

- Dial-path minutes, total and for the top-20 accounts
- GRR bridge: churn vs downsell, by segment
- Volume-variable vs committed-minimum ARR %
- AI-agent ARR, containment %, and gross margin on AI minutes
- Competitive losses to named AI-natives

DILIGENCE ASK

Ask for the automation-rate cohort data and the AI-SKU attach/ARPU bridge. Those two numbers tell you if the AI is real or just a story.

What a credible AI plan must contain

Test management's roadmap against this bar on the call. A credible plan doesn't bolt bots onto human calls; it takes one workflow, makes the AI run it end to end, keeps it on TCN's phone lines, then expands:

Now · 0-6 mo

Pick one workflow, own the line

One high-volume, scriptable workflow (first-party reminders, inbound payment, promise-to-pay). A forward-deployed team to build it. Defaulted onto TCN's own telephony so the AI call stays on TCN's pipes. Containment, cost-per-resolution and liquidation tracked from day one.

Next · 6-18 mo

Run it end to end, price on outcomes

That workflow fully autonomous with compliance in the loop: deterministic pre-call eligibility gating, template-constrained amounts and disclosures, escalation triggers, 100% post-call audit. Priced per resolution. Containment and recovery numbers published with named customers. The phone line locked into the contract.

Later · 18-36 mo

Make it most of the volume

Automated resolutions become the majority of calls, with the outcome data compounding (the flywheel). Expansion workflow by workflow, then into healthcare RCM. Defended with compliance and the audit record. Build isn't the only path: partnering with or acquiring a collections-AI player, or opening Operator as the compliance/telephony substrate for third-party agents, gets there with less capital.

THE TEST

It's not "do you have AI features." It's "can the AI run the call end to end and keep the phone line on TCN." That's the whole game, and it's exactly what the AI-natives are already doing.

Underwrite the calling business, not the AI story

The deal team's "vertical Five9" framing is right: same CCaaS economics in a narrower, more defensible vertical. Here's what that's worth to a lender.

Like Five9

Cloud dialer + omnichannel CCaaS, usage-exposed revenue, AI sold as an add-on, and the same exposure to AI-native disruption of the seat model.

Edge

Collections/compliance depth. Consumption pricing, so AI volume adds revenue instead of killing seats. Last independent specialist after LiveVox went to NICE. Bootstrapped and profitable.

Disadvantage

~14-39x smaller than the scaled leaders, so far less AI R&D. And it faces the same agentic disruption with less capital to fight it.

- **The cash flows that service the debt** come from the vertical and switching costs, not the AI.
- **Don't credit the AI.** Zero agentic revenue through 2025; the \$174M plan rests on a flat price curve.
- **If BSP lends:** take quarterly reporting on dial-path minutes, the GRR bridge, and AI-native losses in the base (see monitoring slide).
- **Risks:** sub-scale R&D, price compression on voice AI, incumbents bundling agents + compliance, AI-natives taking the call.

BOTTOM LINE

Good niche calling business with a modest, real moat from the vertical, switching costs and the pricing model. No AI moat today, and behind on AI, which makes AI a monitorable headwind on this credit rather than upside. Underwrite the existing business; the agentic plan is slideware until management demonstrates it live.

Management-call topics to validate the moat

Proposed topics list for Union Square Advisors, TCN's sell-side advisor. Designed to confirm whether the moat is real, whether AI is traction or narrative, and whether TCN can build a durable moat.

1 · Revenue durability

GRR is 81% while switching costs are "deep". What drives the ~19% gross churn/downsell, and in which segment? How much ARR is volume-variable vs. committed minimums (downside in a soft quarter)?

2 · AI traction, not narrative

AI Voice Agent: live customer count, ARR, containment/automation rate and cost-per-resolution *today*? Do agent calls run on TCN's telephony, or can customers bring their own? (dial-path retention)

3 · Proprietary AI

"In-house model training for Virtual Agents". Trained model, fine-tune, or prompting on GPT/Gemini? What is actually owned beyond SmartAMD? Realized AI ARPU vs. the "4-10× human-minute" claim?

4 · Competition / disintermediation

How often do Skit.ai / Decagon / Sierra / Convoso appear in deals, and the trend? Any volume lost to an AI-native? Post-LiveVox/NICE displacement. Win-rate vs. Five9 in collections, and how durable?

5 · Compliance depth

Is the consent/audit ledger a court-tested, sellable differentiator or configurable rules like peers. Any case where TCN's records were the legal defense? Cost-to-serve to stay current across 50 states + federal + intl?

6 · Scale, economics & strategy

R&D as % of revenue vs. Five9/NICE and the funded AI-natives; does AI expand or compress gross margin? 3-yr plan and targets to convert seats → agents; capital / M&A / partnership posture to lead substitution?

Sources, method & glossary

Primary sources

TCN (tcn.com): AI, Operator, NLC, voice-analytics, careers, case studies, press releases (2025).
 Competitors: Five9, NICE, Talkdesk newsrooms; CX Today.
 Compliance market: DNC.com / Contact Center Compliance, PossibleNOW, Blacklist Alliance.
 Profiles: Crunchbase, ZoomInfo, PitchBook, Top Workplaces; G2 / Capterra reviews.
 BSP: Franklin Templeton disclosures (2025-26).

Method & evidence for the four hardest claims

- "0 customers / \$0 agentic at end-2025": inferred from CIP p. 54 ("AI Voice Agent: First Customer in April 2026") and p. 55 (\$1M 2026E plan). As of this writing (June 2026): one known customer.
- Leadership/AI-talent: based on CIP org pages and public bios (company site, LinkedIn), searched June 2026; no named AI/ML leader found. Negative evidence; test on the call.
- Skit.ai runs end-to-end collections calls with bundled telephony: vendor site and public case studies, June 2026.
- "Rebuilt in an afternoon": four generative features reproduced offline over a 100-record synthetic dataset; excludes telephony, scale, production compliance. Demo-scale evidence of commodity components, no more.

Glossary

CCaaS. Contact center as a service. **ARM**. Accounts-receivable management (collections). **ASR**. Automatic speech recognition. **NLC**. TCN's Natural Language Compliance. **Agentic AI**: AI that autonomously completes multi-step tasks. **Containment**: % of interactions fully handled by AI with no human. **RPC**: right-party contact. **GRR/NRR**: gross/net revenue retention. **Dial path**: the telephony route a call takes (carrier, numbers, reputation). **AMD**: answering-machine detection. **FDE**: forward-deployed engineering. **STIR/SHAKEN**: caller-ID attestation standards. **TCPA / Reg F**. US telemarketing & debt-collection rules.

Companion artifact: a runnable reconstruction of TCN's AI stack (four components + audit ledger) accompanies this deck and is available on request.