

TCN, Inc.. AI Differentiation & Competitive Moat

Is TCN's AI a real moat, or mostly slides? My independent read for BSP.

PREPARED FOR

Benefit Street Partners

MANDATE

AI / technology advisory

SECTOR

CCaaS · Collections / ARM software

DATE

June 2026

PREPARED BY

K. Slowinski
Independent AI & Technology Advisor

Basis of preparation

Scope & purpose

This material was prepared at the request of Benefit Street Partners ("BSP") to support internal evaluation of TCN, Inc. ("TCN" or the "Company"). It addresses AI/technology differentiation and competitive moat only, and is not a complete business, financial, legal, or tax assessment.

Sources & estimates

Analysis relies on public company materials, competitor disclosures, third-party review corpora, and primary technical reconstruction. Financial and headcount figures are third-party estimates, are unaudited, and are labelled where used. No representation or warranty, express or implied, is given as to accuracy or completeness.

No advice / forward-looking

Nothing herein is investment, legal, or regulatory advice or a recommendation to transact. Forward-looking statements involve risk and may differ materially from outcomes. Recipients should perform independent diligence.

Confidentiality

For the internal use of BSP investment and diligence teams only. Not to be reproduced or distributed, in whole or in part, without prior written consent of the author.

The headline: good calling business, no AI moat

VERDICT

TCN has no durable AI moat. Its AI is at competitive parity, built largely on commodity third-party models, and replicable in days. We could not identify a single AI capability TCN offers that its major competitors do not.

- **No edge.** Agent assist, summaries, analytics, voicebots, auto-QA. All table stakes, and Five9, NICE and Talkdesk ship them too, often more maturely.
- **Commodity stack.** Speech, LLMs and even compliance data are bought, not built. TCN is an integrator, not a model owner.
- **The real moat is elsewhere.** Vertical/compliance depth in collections and switching costs. And it is only **modest**.
- **Behind, not ahead.** The only AI-native moat is labor substitution. And TCN's agentic AI had **0 customers and \$0 ARR at end-2025** (first customer Apr 2026). AI-native players are built differently. Until TCN converts seats to agents, AI is a feature, not a franchise.

RECOMMENDATION. CREDIT LENS

*It's a good calling business. Underwrite the loan on that: switching costs, a sticky-ish base, and **consumption pricing** (AI volume makes them money instead of killing seats). Don't credit the AI story, and treat being behind on AI as a headwind to the credit.*

What would make AI a moat. And how we tested it

A capability is only a moat if a well-funded competitor cannot cheaply replicate it or take the customer. We scored TCN against five tests.

Test 1

Proprietary models / IP

Own models, training pipeline, or technical IP rivals lack.

Test 2

Proprietary data & network effects

Unique data that compounds and improves the product with scale.

Test 3

Workflow lock-in

Switching costs, integrations, data gravity, audit/legal record.

Test 4

Distribution / vertical depth

Channel, brand, and domain depth in a defensible niche.

Test 5. The decisive one

Economic step-change: labor substitution

Does the AI remove the customer's largest cost. Human agents. Rather than merely assist them? This is the only AI capability that re-bases unit economics and creates new lock-in. **We weight it highest.**

Methodology: public materials & competitor disclosures · review corpora (G2/Capterra) · third-party headcount/financial estimates · hands-on reconstruction of TCN's AI components.

TCN, Inc.. Small, profitable, vertically focused

ARR (2025A)

~\$60M

~\$69M Q1'26 · ~37% 2026E

Customers · ARPA

~1,200

~\$60K ARPA (up from \$42K)

Profitability

~20→24%

EBITDA LTM → 2026E, 75% GM, NPS 52

Retention

112%

NRR · 81% GRR · 98% logo

Profile

Founded 1999, St. George, Utah. Flagship platform **TCN Operator** (six modules: omnichannel, compliance/data, analytics, workforce, automation, infrastructure). Verticals: **collections / accounts-receivable management**, BPO, healthcare RCM, utilities; actively expanding UK/EU.

Credit read

A genuinely attractive niche credit: ~\$60M ARR, ~37% 2026E growth, ~20 → 24% EBITDA, clears Rule of 40 (their "Rule of 60" is a 2026E projection), bootstrapped (**\$0 raised since 1999**). But **~17-45x smaller** than the three scaled leaders (Five9/Genesys/NICE). Too small to win an AI R&D arms race. Underwrite the calling business, not the AI story.

The AI toolbox. And what's actually under the hood

COMPONENT	MARKETED AS	UNDER THE HOOD (OUR ASSESSMENT)
Conversational / Voice Analytics	Transcribe, search, flag 100% of calls	Third-party ASR + keyword/intent scoring. commodity
Agent Assist	Real-time coaching, next-best-action, summaries	LLM layer over transcripts. commodity wrapper
Natural Language Compliance (NLC)	"Industry-first" plain-English compliance rules	LLM English → rule parser on a rules engine; value is the audit trail, not the AI
Chatbots / Email automation	Deflect & automate routine inbound	LLM + RAG. commodity
Voicebot / AI IVR	Self-service voice, intent routing	ASR + LLM + TTS pipeline. commodity
Auto-QA / Workforce Opt.	Auto-score every conversation	Scoring rubric over transcripts. table stakes

SIGNAL

TCN's generative AI (transcription, summaries, QA) runs on third-party LLMs (GPT / Gemini / Whisper). Its only **in-production** proprietary model, **SmartAMD**, is a narrow answering-machine detector. Real, but not a generative or agentic moat. No general-purpose model or ML research function is evident.


Nothing in the toolbox a major peer doesn't already have

AI CAPABILITY	TCN	FIVE9 (GENIUS)	NICE (ENLIGHTEN/MPower)	TALKDESK (ASCEND)	TCN EDGE?
Real-time agent assist / copilot	●	●	●	●	None
Automated call summaries	●	●	●	●	None
Conversational / voice analytics	●	●	●	●	None
Self-service voicebot / AI IVR	●	●	●	●	None
Auto-QA (100% of calls)	●	●	●	●	None
No-code AI agent builder	○	●	●	●	Behind
Autonomous / agentic AI (acts end-to-end)	◐	●	●	●	Behind
Proprietary models / data moat	◐	◐	●	◐	None
Compliance / TCPA depth (collections)	●	◐	●	◐	Vertical, not AI

● shipping ◐ partial / emerging ○ absent or trailing. TCN's only relative strength is collections-compliance depth. A **domain** advantage, not an AI one. Its only in-production proprietary model (SmartAMD) is narrow answering-machine detection. On true agentic AI it trails the leaders.

A niche specialist in a field of giants

VENDOR	EST. REVENUE / ARR	STRUCTURE	AI SUITE	STRATEGIC FOCUS
NICE	~\$2.7B	Public	Enlighten / Mpower	Enterprise CX, WFO, compliance
Genesys	~\$2.3B	Private	Cloud CX AI	Enterprise omnichannel at scale
Five9	~\$1.0B	Public	Genius AI	Outbound / blended CX, SMB-ent.
Talkdesk	~\$0.2B	Private	Ascend AI	Mid-market / enterprise, AI-first
TCN	~\$60M	Bootstrapped	"Operator" AI	Collections / ARM vertical



IMPLICATION

TCN is ~17-45× *smaller* than the three scaled leaders (Five9/Genesys/NICE); only AI-first Talkdesk (~\$0.2B) is within ~3×. It can't out-spend the majors on AI R&D. It wins by *vertical depth and price*. A defensible niche, not an arms-race position.

Compliance is table stakes, not a moat

The CIP and the commercial VDD both lean on a "compliance moat." But you cannot sell into collections without compliance. And every competitor, including the AI-native disruptors, markets it.

VENDOR	MARKETS COLLECTIONS COMPLIANCE?	WHAT THEY CLAIM
LiveVox (now NICE)	● Yes. Its heritage	Pioneered "Four Clouds" / Human Call Initiator; court-validated non-ATDS
Five9	● Yes	Dedicated TCPA and debt-collections product pages
Convoso	● Yes	Compliance product: dialing guardrails + automation in-workflow
Skit.ai (AI-native)	● Yes. Headline	"Automatically follows FDCPA, TCPA, Reg F". TCN's NLC pitch, verbatim
NICE / Genesys / Talkdesk	● Yes	Compliance tooling, audit trails, Reg-F templates
DNC.com / Gryphon / PossibleNOW	● Yes	Compliance-as-a-service. Buyable a la carte

WHY IT MATTERS

Compliance marketing is table stakes; compliance depth (court-tested audit records, a revocation ledger) is a modest, real **vertical/legal** advantage. But not an AI one, and eroding as the AI-natives go compliance-native. Strip the label and the "moat" is switching costs + being the last independent specialist after LiveVox was absorbed by NICE.

The AI is rented; the moat would have to be built

Bought / commodity

- ASR. Whisper / Deepgram / AWS Transcribe class
- LLMs. OpenAI / Anthropic / AWS
- TTS. ElevenLabs / Cartesia / Polly class
- **Compliance data.** DNC, litigator & reassigned-number lists are a commodity market (DNC.com: 70B+ scrubs; PossibleNOW; Gryphon; Blacklist Alliance)


Built / defensible

- Integrated platform & dialer
- Contact-eligibility rules engine + NL authoring
- **Consent / revocation audit ledger** (the courtroom record)
- Collections-specific workflows & data gravity

REPRODUCIBILITY TEST. THE IP BARRIER IS LOW

*As part of this diligence we reproduced TCN's four headline **AI/LLM components** offline over a 100-record dataset in an afternoon. Confirming the AI layer is commodity. What we did **not** reproduce is the production consent/audit ledger at customer scale, with years of legally-defensible history. Which is exactly where the modest, non-AI moat sits.*

Where the real defensibility is (it isn't AI)



What's real

Purpose-built for debt collection (a customer literally cites "solutions designed for the debt-collection industry"), reinforced by switching costs and the legal-defensibility of consent/audit records.

What's not

The AI. It neither creates lock-in beyond the platform itself nor compounds with proprietary data. Removing the "AI" branding would not change TCN's competitive position.

The only AI moat is labor substitution

Contact-center cost stack

~60-70%

is human agent labor. The prize

Augmentation (today)

Copilots & summaries trim handle time. **Incremental**.
And equally available to every competitor. No lock-in.

Substitution (the moat)

Replace seats with autonomous agents → re-bases unit economics, shifts pricing seat → outcome, and creates an automated-resolution **data flywheel** + deep lock-in.

THESIS

*A moat is created only by convincing customers to remove human agents and replace them with AI agents. Industry-wide that conversion has **not** happened at scale. "agentic" is still largely demoware. The first mover to convert seats into agents captures the labor budget and the lock-in. **Until TCN does this, there is no moat.***

TCN-specific angle

Collections is unusually automatable. Scripted, repetitive, compliance-bounded. So the **opportunity is genuine**. But TCN sells augmentation today and is sub-scale on R&D. This is the single most important item to diligence and to fund.

Who's actually coming for the labor budget

The \$14.4B "AI agent opportunity" the CIP/VDD book as TCN upside is the same prize these well-capitalized AI-natives are racing for. **Neither document names a single one of them.**

COMPANY	WHAT IT DOES	SCALE	THREAT VECTOR TO TCN
Decagon	Autonomous "concierge" AI agents resolving chat/email/voice to cut support headcount; plain-English auditable compliance (AOPs) + QA guardrails	\$4.5B val · ~\$35M ARR (est. late-25) · 2023	Highest. Launching outbound voice + NLC-style compliance
Sierra	Autonomous customer-facing AI agents (voice + digital), outcome-based pricing; founded by Bret Taylor (ex-Salesforce / OpenAI chair)	\$15.8B val · ~\$200M ARR · 2023	High. Capital + trust to move into regulated
Skit.ai	Collections-native voice AI: end-to-end outbound collections calls on its own telephony (owns the dial path), with built-in FDCPA / TCPA / Reg F compliance	~\$48M raised · 2016	Direct. TCN's exact use case & buyer
Cresta	Generative AI for the contact center: real-time agent assist, coaching, and virtual agents; analytics-driven	\$1.6B val · ~\$100M ARR · 2017	Overlaps agent-assist / analytics; FS vertical
Parloa	Voice-first agentic AI automating inbound/outbound conversations at enterprise scale	\$3.0B val · 2018	Voice automation; EU/enterprise-led

TCN'S DEFENSES: HOLDING VS. ERODING

Holds: consumption pricing, incumbency in the dial path, vertical/compliance depth, switching costs. **Eroding:** compliance is now native to the disruptors, and the AI-natives can **become the dial path**. Skit.ai already offers its own telephony and end-to-end collections calls today. "Watch them" is the floor. The real question is dial-path retention, and whether TCN partners, acquires, or builds fast enough.

What they're already doing (and TCN isn't)

Same prize everyone points at: the labor budget. These are the folks actually going after it, faster, with far more money, and several already own the phone line.

COMPANY	WHAT THEIR AI DOES	END TO END?	OWNS THE LINE?	MONEY	VOICE / COLLECTIONS TODAY
TCN	Assists the human agent (summaries, QA, answering-machine detection)	○ no	● its dialer	Bootstrapped, ~\$60M ARR	Behind: 0 AI-agent users end-2025
Five9 / NICE	Copilots plus early agents	◐ partial	● their dialer	Public, ~\$1-2.7B rev	Incumbent, not collections-first
Decagon	Resolves calls end to end, cuts headcount	● yes	◐ entering voice	\$4.5B	Enterprise, moving to outbound voice
Sierra	Autonomous agents, voice and digital	● yes	● voice	\$15.8B, ~\$200M ARR	Broad, regulated next
Skit.ai	Runs the whole collections call	● yes	● own telephony	~\$48M raised	Collections-native, live now
Intercom (Fin)	Fin Voice: real-time phone workflows, outbound	● yes	● voice	Public	Voice live, expanding

READ IT

These players ship autonomous, end-to-end calls today, carry 25-250x TCN's capital, and several **already own the phone line**. TCN assists humans and booked **zero** AI-agent revenue last year. And note their own CIP prices AI voice at **\$0.20-0.50/min** vs **~\$0.05 human**: the whole AI plan is trading phone minutes for AI minutes, which only works if the call stays on TCN's line.

TCN builds AI for the human, not the agent

What TCN does

Adds AI to the human's workflow: agent assist, call summaries, QA scoring, better answering-machine detection. Useful. But a human is still on the call. It's incremental.

What the AI-natives do

Rethink the whole process. The AI runs the call end to end, no human, starting with one narrow workflow and expanding. That's how you actually take the labor budget.

TWO PROBLEMS

Labor dollars: TCN keeps saying it'll win them. It hasn't picked up any yet. **Talent:** leadership is the Bird family (Terrel, Jesse, Darrin) plus a 30-year software vet, all deep contact-center. I don't see AI/ML research leadership or named AI hires. The CIP claims "AI talent." Where is it? If they were AI-native they wouldn't be this far behind.

The call

They've got a good platform. Double down, AI-enable it, and move a lot faster. Right now they're behind, and the people who aren't behind can walk into their accounts.

They measure a better dialer, not better collections

TCN can quantify a better dialer and cheaper QA. It cannot yet quantify the metrics that actually decide collections. The ones the AI-natives lead with.

METRIC THAT MATTERS TO A COLLECTIONS STAKEHOLDER	WHAT TCN CAN SHOW TODAY	THE AI-NATIVE BAR (SKIT.AI / PRODIGAL)
Right-party contact / connect	● RPC 7.68% → 8.70%, +13% agent utilization, 8.8× ROI (SmartAMD)	RPC tracked across 1B+ conversations, \$1B+ resolved
QA / compliance coverage	● 100% of calls AI-scored, 50% fewer false positives, 30% infra cut	FDCPA/TCPA/Reg F enforced natively in the agent
Autonomous resolution / containment %	. Zero through end-2025: 0 customers / \$0 (first customer Apr 2026)	AI handles high-volume follow-ups; escalates only complex cases
Cost per resolution vs. human	. Not reported	Tracked per client; 10× ROI case studies
Liquidation / recovery lift from AI	. Not reported (connect-rate only)	Intent-to-pay scoring lifts liquidation, cuts collection time (Prodigal)

BOTTOM LINE

*Their numbers prove a good dialer and cheaper QA. That's efficiency. This category gets won on **containment, cost-per-resolution and liquidation**, and TCN can't report one of those yet. Strip the AI label and the numbers say: a very good calling platform.*

Scored against the five tests

MOAT TEST	RATING	RATIONALE
1 · Proprietary models / IP	LOW	One narrow proprietary model (SmartAMD, AMD); generative AI is third-party. No franchise-grade IP.
2 · Proprietary data / network effects	LOW	Compliance data is bought; no compounding data advantage.
3 · Workflow lock-in / switching costs	MEDIUM	Real via integration + consent/audit records; not AI-driven.
4 · Distribution / vertical depth	MED-HIGH	Genuine collections/ARM domain depth; efficient GTM.
5 · Labor substitution (weighted x)	NOT YET	Sells augmentation, not replacement. The decisive, unrealized lever.
Overall AI moat	LOW	Niche franchise with modest, non-AI defensibility. No AI premium warranted.

Ratings are qualitative judgments for diligence framing, not a quantitative score.

KPIs & catalysts to monitor

Substitution is happening (bull signals)

- % interactions **fully automated** / containment rate rising
- Reference customers **removing agent seats**
- Pricing shifts from **per-seat** → **per-outcome**
- **Gross margin expansion** from automation
- Agentic product GA with named logos

Still a feature (bear signals)

- AI sold only as add-on; **low SKU attach / ARPU uplift**
- **R&D % of revenue** trails peers (scale gap widens)
- Logo churn in collections vertical
- Larger peers bundle **agentic + compliance** and undercut
- Net revenue retention flat/declining
- AI agents need **forward-deployed engineering**. Services-gated, not self-serve

DILIGENCE ASK

Ask for the automation-rate cohort data and the AI-SKU attach/ARPU bridge. Those two numbers tell you if the AI is real or just a story.

If I ran AI at TCN, here's what I'd build

Kathy Slowinski, if I ran AI/product here: I wouldn't bolt bots onto human calls. I'd take one workflow, make the AI run it end to end, keep it on TCN's phone lines, then expand. Roughly:

Now · 0-6 mo

Pick one workflow, own the line

Take one high-volume, scriptable workflow (first-party reminders, inbound payment, promise-to-pay). Stand up a forward-deployed team to build it. Default it onto TCN's own telephony so the AI call stays on TCN's pipes. Track containment, cost-per-resolution and liquidation from day one.

Next · 6-18 mo

Run it end to end, price on outcomes

Make that workflow fully autonomous, no human, compliance in the loop. Charge per resolution, not per seat. Publish containment and recovery numbers with named customers. Lock the phone line into the contract.

Later · 18-36 mo

Make it most of the volume

Automated resolutions become the majority of calls. Use that data to keep getting better (the flywheel). Expand workflow by workflow, then into healthcare RCM. Defend with compliance and the audit record.

THE TEST

It's not "do you have AI features." It's "can the AI run the call end to end and keep the phone line on TCN." That's the whole game, and it's exactly what the AI-natives are already doing.

Underwrite the calling business, not the AI story

- **The cash flows are real.** They come from the vertical and switching costs, not the AI. That's what services the debt.
- **Don't credit the AI.** It's slideware today, and being this far behind makes AI a headwind to the credit, not upside.
- **Downside's contained** by sticky-ish collections workflows and the audit records. Good for a credit.
- **If we lend:** monitor dial-path volume and AI-native displacement in the base, and push them to run one workflow end to end and keep the phone line. That's their only path to a real moat.
- **Risks:** too small on R&D, CCaaS going commodity, big players bundling agents + compliance, and the AI-natives taking the call.

BOTTOM LINE

*Good niche calling business, modest moat from the vertical and switching costs, **no AI moat today**. The AI only matters if they actually replace human agents, and they haven't started. Price what's real. Treat the AI as an option, not the reason to buy.*

Is there something here? Yes, a "vertical Five9"

Leon's framing is right. TCN is a collections/ARM Five9: same CCaaS, usage-based economics, in a narrower and more defensible vertical. The question is what that's worth.

Like Five9

Cloud dialer + omnichannel CCaaS. Usage-exposed revenue. AI sold as an add-on. And like Five9, exposed to AI-native disruption of the seat model.

TCN's edge vs. Five9

Collections/compliance depth. **Consumption** pricing, not per-seat, so AI adds revenue instead of killing seats. Last independent standing after LiveVox went to NICE. Bootstrapped and profitable.

TCN's disadvantage

~17x smaller than Five9, so far less AI R&D. Thinner enterprise footprint. And a "vertical Five9" faces the same agentic disruption with less money to fight it.

VERDICT

*Worth pursuing. A real, defensible **niche business** (clears Rule of 40, sticky-ish, consumption-priced). But the moat is **the vertical, switching costs and the pricing model, not AI**. Underwrite it as a lender on the existing business; the future business is slideware. The AI headwind is the thing the management call has to test.*

Management-call topics to validate the moat

Proposed topics list for Union Square. Designed to confirm whether the moat is real, whether AI is traction or narrative, and whether TCN can build a durable moat.

1 · Revenue durability

GRR is 81% while switching costs are "deep". What drives the ~19% gross churn/downsell, and in which segment? How much ARR is volume-variable vs. committed minimums (downside in a soft quarter)?

2 · AI traction, not narrative

AI Voice Agent: live customer count, ARR, containment/automation rate and cost-per-resolution *today*? Do agent calls run on TCN's telephony, or can customers bring their own? (dial-path retention)

3 · Proprietary AI

"In-house model training for Virtual Agents". Trained model, fine-tune, or prompting on GPT/Gemini? What is actually owned beyond SmartAMD? Realized AI ARPU vs. the "4-10× human-minute" claim?

4 · Competition / disintermediation

How often do Skit.ai / Decagon / Sierra / Convoso appear in deals, and the trend? Any volume lost to an AI-native? Post-LiveVox/NICE displacement. Win-rate vs. Five9 in collections, and how durable?

5 · Compliance depth

Is the consent/audit ledger a court-tested, sellable differentiator or configurable rules like peers. Any case where TCN's records were the legal defense? Cost-to-serve to stay current across 50 states + federal + intl?

6 · Scale, economics & strategy

R&D as % of revenue vs. Five9/NICE and the funded AI-natives; does AI expand or compress gross margin? 3-yr plan and targets to convert seats → agents; capital / M&A / partnership posture to lead substitution?

Sources, method & glossary

Primary sources

TCN (tcn.com): AI, Operator, NLC, voice-analytics, careers, case studies, press releases (2025).
 Competitors: Five9, NICE, Talkdesk newsrooms; CX Today.
 Compliance market: DNC.com / Contact Center Compliance, PossibleNOW, Blacklist Alliance.
 Profiles: Crunchbase, ZoomInfo, PitchBook, Top Workplaces; G2 / Capterra reviews.
 BSP: Franklin Templeton disclosures (2025-26).

Method

Triangulation of public materials, competitor disclosures and review corpora; third-party financial/headcount estimates (unaudited); and a primary technical reconstruction of TCN's AI components incl. a tamper-evident consent/audit ledger over 100 records. Ratings are qualitative diligence framing.

Glossary

CCaaS. Contact center as a service. **ARM**. Accounts-receivable management (collections). **ASR**. Automatic speech recognition. **NLC**. TCN's Natural Language Compliance. **Agentic AI**. AI that autonomously completes multi-step tasks. **TCPA / Reg F**. US telemarketing & debt-collection rules.

Companion artifact: a runnable reconstruction of TCN's AI stack (four components + audit ledger) accompanies this deck and is available on request.