

Call Prep: What Good and Bad Answers Sound Like

Companion to the 5-item agenda. One goal per item; keep them off slides and on specifics.

1 Retention quality: is the 81% GRR a compression event?

Good: a clean bridge; downsell concentrated in small accounts or a one-time repricing; volume-variable share modest with minimums underneath; downsell cohorts stable or improving.

Bad: can't produce the bridge; downsell spread across the base or concentrated in large accounts; "usage variability" with no cause; any account shrinking because volume moved to another vendor's automation.

Follow-up: "Name the three largest downsells last year and why." If any answer involves a customer automating with someone else, that's the displacement thesis showing up in the numbers.

2 Pricing: efficiency pricing vs outcome pricing

Good: they acknowledge per-minute AI pricing will compress, have modeled it, and have an outcome-based (per-resolution) pilot or roadmap; they can articulate what they still monetize if the agent isn't theirs (compliance, numbers, recording, analytics).

Bad: they defend \$0.30/min flat for five years; no outcome-pricing plans; no answer for the bring-your-own-agent scenario beyond "customers won't do that."

Follow-up: "The winning AI companies price on outcomes. When do you?" Their reaction tells you if management sees the category shift.

3 Roadmap credibility: dates, owners, inventory

Good: named owner for AI, dated milestones, honest shipped-vs-planned split; a crisp answer on the native-stack denominator; an inbound-first automation plan (FCC limits outbound autonomy, inbound conversion is the near-term prize).

Bad: "expanding portfolio of proprietary models" with no names; in-development items with no dates or owners; conflating SmartAMD execution volume with generative capability; outbound-first agentic promises that ignore the FCC posture.

Follow-up: "Measured by generative capability, not per-call executions, what fraction of your AI runs on GPT, Gemini, Mistral or Whisper?" There is no good evasion of this one.

4 Implementation: who actually stands this up?

Good: a real services answer: named team, deployments measured in weeks, cost per deployment declining across cohorts, a repeatable playbook for mid-market accounts that can't self-configure.

Bad: "it's self-serve" or "deploy in a day" for AI agents; no forward-deployed capability and no plan to build or partner for one. The CIP has no professional-services story at all; if the call doesn't either, agentic revenue has no delivery mechanism.

Follow-up: "Walk me through the April customer's deployment: who did the work, how long, what did it cost you?"

5 The live demo: the whole call in one moment

Good: a live call, end to end: disclosures, a negotiation turn, a clean human handoff; real containment and cost-per-resolution numbers from the live customer, even if modest.

Bad: a video, a slide, "our lawyers won't allow live," or a demo that dodges the compliance disclosures. A polished recording is a produced artifact; treat refusal to go live as a data point in itself.

Follow-up if they refuse: "Then give us read access to a sandbox and we'll place the call ourselves."

REMINDERS FOR THE ROOM

- **Watch list, not agenda:** if Sierra or Decagon is acquired by Five9, Talkdesk or Genesys, the threat level steps up; that's a monitoring trigger for the loan, not a management question.
- **Number check:** BSP cited Sierra at "~\$400M raised, ~\$4B valuation" on the prep call; current public marks are ~\$950M raised at \$15.8B (May 2026). Decagon is the \$4.5B one. Worth quietly correcting so nobody anchors low on the threat.
- **Don't accept averages:** every retention and containment answer should come by cohort or by named account, not blended.
- **Timing:** push Union Square to lock the slot this week; sell-side delay favors the seller.