

Management Call — Diligence Topics & Questions

Validating TCN's technology moat and other sources of defensibility

PREPARED FOR **Benefit Street Partners — to be shared with Union Square Advisors**
 PREPARED BY **K. Slowinski — Independent AI & Technology Advisor**
 DATE **June 2026**

OBJECTIVE

TCN presents as a "vertical Five9" for collections/ARM. The business is genuinely attractive (~\$60M ARR, Rule-of-60, sticky, consumption-priced), but the CIP and commercial VDD assert an "**AI + compliance moat**" that the underlying evidence does not yet support. These questions are designed to establish (i) whether the technology/AI moat is real or narrative, (ii) which non-AI moats are durable, and (iii) whether management has a credible plan to *build* a moat as AI reshapes the category.

PART A — TECHNOLOGY & AI MOAT (PRIORITY)

A1 · Proprietary models & IP

Testing: is any AI genuinely owned, or a wrapper over third-party LLMs?

1. For each core AI feature (transcription, summarization, sentiment, QA, virtual agents), which specific vendor/model powers it (GPT / Gemini / Whisper / Mistral), and which are owned in-house?
2. "In-house model training to enhance Virtual Agents" — is this a proprietary trained model, fine-tuning/LoRA on a base model, or prompt/RAG engineering? What exactly is the owned IP?
3. SmartAMD: architecture, training-data source and volume, retrain cadence, and measured accuracy vs. standard AMD and vs. competitors. Is the 8.8x ROI repeatable across customers, or a single case?
4. What proprietary training data do you hold that a competitor cannot obtain or buy? Do models measurably improve with usage (a data flywheel) — can you show the curve?
5. As frontier LLMs commoditize transcription/summarization/agents, where does durable technical advantage remain?

A2 · AI agents — traction vs. narrative

Testing: is agentic real revenue or slides? (First customer April 2026 — i.e., 0 customers / \$0 at end-2025.)

1. **Live demonstration — not slides.** Show the AI Voice Agent handling an end-to-end collections call in real time (compliance disclosure, negotiation / promise-to-pay, human escalation), plus the live customer's usage/metrics dashboard.
2. AI Voice Agent: live production customers today, ARR, and go-live dates — and what were these figures at 12/31/2025?
3. For live deployments: containment / full-automation rate (% of interactions with no human), cost-per-resolution vs. human, and any liquidation/recovery or PTP lift.
4. Which collections use cases are automated (inbound payment, promise-to-pay, first-party reminders, disputes) and which are off-limits, and why?
5. 2026–27 AI-agent pipeline, conversion assumptions, and the build behind the \$1M → scale ramp.
6. Compliance controls on autonomous calls (mini-Miranda, revocation handling, recording) and the human-escalation design.
7. What does it take to stand up an AI agent for one customer — hours of TCN forward-deployed / professional-services work, time-to-live, and cost — and is it self-serve or bespoke? How does that scale across a mostly mid-market base of 1,000+ accounts?

A3 · Dial-path retention / disintermediation (the crux)

Testing: does TCN stay in the value chain as AI takes over the conversation?

1. Do AI-agent interactions run on TCN's own telephony/carrier stack, or can a customer point a third-party agent at its own SIP and bypass TCN? The AI-natives **already offer their own telephony** (e.g., Skit.ai runs end-to-end collections calls today) — what specifically keeps the dial path with TCN?
2. What technically or contractually keeps the dial path with TCN as volume shifts from human to AI?
3. Have you lost any call volume or a customer to an AI-native agent? Is any customer running a third-party agent alongside TCN today?
4. If a customer brings its own AI agent, what does TCN still monetize (carrier, compliance, recording, analytics), and at what margin?

A4 · Architecture, scale & AI unit economics

Testing: can the platform and the margins support AI at scale?

1. Gross-margin impact of AI: LLM/inference + carrier cost per AI minute vs. price charged — does AI *expand* or *compress* gross margin as it scales? Show per-conversation unit economics.
2. Real-time voice latency, multi-tenant limits, uptime/SLAs for autonomous voice.
3. Dependency risk on third-party model vendors and orchestration frameworks (LangChain/MCP) — exposure if a vendor changes pricing or terms.

A5 · AI roadmap & the plan to build a moat

Testing: is there a credible path from feature to franchise?

1. Three-year plan to convert seats → agents: targets for automation rate, AI ARR, and AI ARPU.
2. What will you build that a competitor cannot replicate within 12 months — and where does the durability come from (data, workflow, compliance-in-the-loop, distribution)?
3. Of the 120+ engineers, how many are on AI/ML specifically, and what is the hiring plan?
4. Capital, M&A, or partnership posture to lead agentic substitution against AI-natives funded at \$1.5–15B.

PART B — OTHER MOATS & THE BUSINESS

B1 · Switching costs & retention quality

Testing: how sticky is the base, really?

1. Reconcile ~98% logo retention with 81% GRR — what is the gross \$ churn vs. downsell, by segment/cohort, and why does downsell occur under a consumption model?
2. Mechanically, what is the switching cost — integrations, data/records migration, compliance/audit history, retraining? Time and cost for a customer to leave.
3. NRR 112%: expansion drivers (module cross-sell vs. volume vs. price) and their durability.
4. Logo and revenue concentration — top 10 / top 20 as % of ARR.

B2 · Consumption pricing & monetization

Testing: is usage pricing a strength or a hidden volume risk?

1. Share of ARR that is volume-variable vs. committed minimums / platform fees; downside in a soft collections quarter.
2. Seasonality and macro sensitivity of volume to the delinquency cycle.
3. Net price trends per minute/message; discounting under competition; durability of AI-minute pricing vs. the "4–10× human minute" claim.

B3 · Compliance moat — depth vs. marketing

Testing: is compliance a defensible edge, or table stakes everyone markets?

1. Is your compliance legally defensible in litigation — any case where TCN's consent/audit records were the successful defense (analogous to LiveVox's court-tested non-ATDS status)?
2. Is the consent/revocation audit ledger a productized, tamper-evident, exportable capability, or internal plumbing?
3. Coverage and maintenance across 50 states + federal + international (Ofcom, CRTC): team size, update process, cost-to-serve, and speed to implement a rule change.
4. How does your compliance concretely differentiate from Five9 / NICE / Convoso and the compliance-native AI startups?

B4 · Distribution, GTM & vertical position

Testing: durability of the go-to-market advantage.

1. Win/loss and win-rate vs. Five9, LiveVox/NICE, Alvaria, and the AI-natives, by segment, over the last eight quarters.
2. CAC, payback, and the sales motion (inbound vs. outbound vs. partner/referral); sales & marketing as % of revenue.
3. Post-LiveVox/NICE displacement — number of wins and how durable the tailwind is.
4. International (UK grew 9× in 9 months): CAC/payback, compliance fit, and durability.

B5 · Scale, financials & capital

Testing: can a sub-scale, bootstrapped player sustain and extend the moat?

1. R&D as % of revenue vs. peers; total headcount mix (engineering vs. sales vs. services).
2. 2026E reacceleration to ~37% (from ~20% in 2025): driven by price, volume, new logos, or a few large deals? Pipeline coverage of the \$83M target and back-half weighting.
3. Durability of Rule-of-60 and the path to ~45% EBITDA by 2030 — what has to be true.
4. Capital needs and appetite (bootstrapped to date) and any use of proceeds.

B6 · Management, organization & execution

Testing: can this team execute an AI pivot?

1. Founder-led — depth of bench, succession, and single-person dependency.
2. Professional-services / implementation dependency — is "deploy in a day" real at scale, or is there a services bottleneck?
3. AI agents require forward-deployed engineering (the Sierra/Decagon norm: \$50–200K implementation, 8–16 weeks to live). Do you have that capability in-house today, and what does it cost to serve a mid-market collections buyer who cannot self-configure?
4. Key-person and attrition risk in engineering / AI.

WHAT WOULD CHANGE OUR VIEW

Upgrades the thesis

Real containment / cost-per-resolution data on live AI agents; contractual or technical dial-path lock-in; genuinely owned, trained models beyond SmartAMD; a demonstrable data flywheel; agentic ARR ramping with named references.

Downgrades the thesis

Customers can bring their own telephony/agent; GRR erosion or heavy downsell; volume already lost to an AI-native; gross-margin compression from AI; the "AI moat" resting on third-party LLMs plus SmartAMD alone.